关于模型优化做的工作

完成DCU移植

libtorch提供C++接口，性能比Pytho版本提升0.1（有总比没有强）

在此切换Torch_DIR

```
export LD_LIBRARY_PATH="/home/worldpeace/soft/libtorch/lib:$LD_LIBRARY_PATH"
export PATH="/home/worldpeace/soft/deepmd-c++/bin:$PATH"
export LD_LIBRARY_PATH="/home/worldpeace/soft/deepmd-c++/lib:$LD_LIBRARY_PATH"
cmake -L -C ../cmake/presets/basic.cmake  \
-C ../cmake/presets/kokkos-openmp.cmake \
-C ../cmake/presets/kokkos-cuda.cmake \
-DCMAKE_BUILD_TYPE=Release \
#   取消注释二选一    #
#原始版本# -DTorch_DIR=`python -c 'import
torch;print(torch.utils.cmake_prefix_path)'`/Torch \
#libtorch版本# -DTorch_DIR=/home/worldpeace/soft/libtorch/share/cmake/Torch \
-DGFLAGS_INCLUDE_DIR=/home/worldpeace/soft/libtorch/include \
-DCUDA_ARCH=AMPERE86 \
-DMKL_INCLUDE_DIR=/opt/intel/oneapi/mkl/latest/include \
-DCMAKE_PREFIX_PATH=/home/worldpeace/soft/deepmd-kit\
-DCMAKE_INSTALL_PREFIX=/opt/LMP_dp_allegro_C -DBUILD_TOOLS=ON -
DBUILD_SHARED_LIBS=ON \
-DPKG_GPU=ON  \
-DFFT=FFTW3 -DFFTW3_LIBRARY=/opt/fftw3/lib/libfftw3.so \
-DFFTW3_INCLUDE_DIR=/opt/fftw3/include \
-DLAMMPS_INSTALL_RPATH=ON  ../cmake
```

TVM调研实践，卡在Input

input代码分析，获得Input形状

/home/worldpeace/anaconda3/envs/tvm/lib/python3.11/site-packages/nequip/ase/nequip_calculator.py

```
    def calculate(self, atoms=None, properties=["energy"],
system_changes=all_changes):
        """
        Calculate properties.

        :param atoms: ase.Atoms object
        :param properties: [str], properties to be computed, used by ASE
internally
        :param system_changes: [str], system changes since last calculation, used
by ASE internally
        :return:
        """
        # call to base-class to set atoms attribute
        Calculator.calculate(self, atoms)
```

```
        # prepare data
        data = AtomicData.from_ase(atoms=atoms, r_max=self.r_max)
        for k in AtomicDataDict.ALL_ENERGY_KEYS:
            if k in data:
                del data[k]
        data = self.transform(data)
        data = data.to(self.device)
        data = AtomicData.to_AtomicDataDict(data)

        # predict + extract data
        out = self.model(data)
```

这里data为模型需要的输入，获得并保存在txt内部

```
{'edge_index': tensor([[ 0,  0,  0,  ..., 70, 70, 70],        [ 3,  5,  6,  ...,
47, 42, 18]], device='cuda:0'), 'pos': tensor([[1.1242e-01, 1.2245e+01,
4.1290e+00],        [3.6315e-01, 4.2438e+00, 8.0531e+00],        [9.3500e-02,
4.5374e+00, 4.3283e+00],        [1.2334e-01, 8.6023e+00, 8.1177e+00],
[1.2322e+01, 8.6604e+00, 4.3235e+00],        [4.2572e+00, 1.2474e+01,
8.0820e+00],        [4.3863e+00, 1.2205e+01, 4.0424e+00],        [4.3404e+00,
3.7507e+00, 8.1078e+00],        [4.3428e+00, 3.9973e+00, 3.9187e+00],
```

添加tensordict库，修改格式为可识别

```
input=AtomicData.to_AtomicDataDict({'edge_index': tensor([[ 0,  0,  0,  0,  0,
0,  0,  0,  0,  0,  0,  0,  0,  0,          0,  0,  0,  0,  0,
0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,          0,  0,  0,  0,  0,
0,  0,  0,  0,  0,  0,  0,  1,  1,  1,  1,  1,  1,          1,  1,  1,  1,  1,
1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,          1,  1,  1,  1,  1,
1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,          1,  1,  1,  1,  1,
1,  1,  2,  2,  2,  2,  2,  2,  2,  2,  2,  2,  2,          2,  2,  2,  2,  2,
2,  2,  2,  2,  2,  2,  2,  2,  2,  2,  2,  2,  2,          2,  2,  2,  2,  2,
2,  2,  2,  2,  2,  2,  2,  2,  2,
```

jit.load script model。然后运行

```
import torch
from torch import tensor
from nequip.data import AtomicData, AtomicDataDict
from tensordict.tensordict import TensorDict
mod=torch.jit.load('deployed.pth')
mod=mod.to('cuda')
input=AtomicData.to_AtomicDataDict({'edge_index': tensor([[
out=mod(input)
print(out)
```

正确性对比：

上面两个python形式已经对比无误

nequip_calculator.py out 和LAMMPS pair_allegro.cpp的out对比即可

ASE仿真，完成模型独立运行

Script model保存机制，C++调用单步性能插桩

pair_allegro.cpp输出

```
Per MPI rank memory allocation (min/avg/max) = 5.31 | 5.31 | 5.31 Mbytes   Step
       Time           PotEng         KinEng         TotEng           Temp
Press          Volume        Density
0   0             -5115.3514      247.66244     -4867.6889      1000
124912.37      35001.599      9.8101743
model.forward Time is : 0.445529 stoTensor().cpu Time is : 0.000252 s
Pair All Time is : 0.452624 s
model.forward Time is : 3.49675 s
toTensor().cpu Time is : 7.3e-05 s
Pair All Time is : 3.50579 s
model.forward Time is : 0.283893 s
toTensor().cpu Time is : 0.000103 s
Pair All Time is : 0.286683 s
model.forward Time is : 0.282643 s
toTensor().cpu Time is : 6e-05 s
Pair All Time is : 0.290819 s
model.forward Time is : 0.283949 s
toTensor().cpu Time is : 0.000106 s
Pair All Time is : 0.28688 s
model.forward Time is : 0.282495 s
toTensor().cpu Time is : 5.4e-05 s
Pair All Time is : 0.286974 s
model.forward Time is : 0.282829 s
toTensor().cpu Time is : 0.000113 s
```

python 单独load Model对比

```
/home/worldpeace/anaconda3/envs/tvm/lib/python3.11/site-
packages/nequip/__init__.py:20: UserWarning: !! PyTorch version 2.5.1 found.
Upstream issues in PyTorch versions 1.13.* and 2.* have been seen to cause
unusual performance degredations on some CUDA systems that become worse over
time; see https://github.com/mir-group/nequip/discussions/311. The best tested
PyTorch version to use with CUDA devices is 1.11; while using other versions if
you observe this problem, an unexpected lack of this problem, or other strange
behavior, please post in the linked GitHub issue.  warnings.warn(
Time :0.6278s
Time :0.5497s
Time :1.6636s
Time :0.3279s
Time :0.0221s
Time :0.0236s
Time :0.0239s
Time :0.0243s
Time :0.0231s
Time :0.0288s
```

性能分析，不同硬件平台底层走不同计算库

Z100

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA %	CUDA total	CUDA time avg	# of Calls
aten::mm	0.44%	1.953s	0.45%	1.992s	40.903us	90.255s	43.12%	90.337s	1.855ms	48696
aten::_index_put_impl_	18.67%	82.695s	35.08%	155.393s	7.251ms	64.497s	30.81%	74.646s	3.483ms	21432
void (anonymous namespace)::indexing_backward_kernel...	0.00%	0.000us	0.00%	0.000us	0.000us	55.273s	26.40%	55.273s	11.661ms	4740
Cijk_Ailk_Bljk_SB_MT128x64x8_SN_APM1_AF0EM1_AF1EM1_A...	0.00%	0.000us	0.00%	0.000us	0.000us	20.940s	10.00%	20.940s	3.313ms	6320
Cijk_Ailk_Bljk_SB_MT256x32x8_SN_APM1_AF0EM1_AF1EM1_A...	0.00%	0.000us	0.00%	0.000us	0.000us	17.270s	8.25%	17.270s	3.693ms	4676
hipLaunchKernel	0.87%	3.862s	0.88%	3.907s	6.344us	14.560s	6.96%	14.569s	23.657us	615822
Cijk_Alik_Bljk_SB_MT128x128x16_SN_APM1_AF0EM1_AF1EM1...	0.00%	0.000us	0.00%	0.000us	0.000us	14.055s	6.71%	14.055s	3.558ms	3950
aten::mul	0.53%	2.339s	0.55%	2.422s	20.371us	12.128s	5.79%	12.549s	105.547us	118894
Cijk_Ailk_Bjlk_SB_MT64x64x16_SN_APM1_AF0EM1_AF1EM1_A...	0.00%	0.000us	0.00%	0.000us	0.000us	11.667s	5.57%	11.667s	4.990ms	2338
aten::sum	0.11%	486.838ms	0.12%	518.128ms	23.079us	11.580s	5.53%	11.705s	521.388us	22450
void at::native::reduce_kernel<512, 1, at::native::R...	0.00%	0.000us	0.00%	0.000us	0.000us	11.536s	5.51%	11.536s	563.297us	20480
Cijk_Alik_Bljk_SB_MT64x128x16_SN_APM1_AF0EM1_AF1EM1_...	0.00%	0.000us	0.00%	0.000us	0.000us	7.551s	3.61%	7.551s	4.779ms	1580
aten::copy_	0.67%	2.971s	0.70%	3.106s	40.148us	7.321s	3.50%	8.057s	104.132us	77371
void at::native::legacy::elementwise_kernel<128, 4, ...	0.00%	0.000us	0.00%	0.000us	0.000us	6.801s	3.25%	6.801s	229.304us	29660
void at::native::modern::elementwise_kernel<at::nati...	0.00%	0.000us	0.00%	0.000us	0.000us	5.709s	2.73%	5.709s	78.041us	73152
void (anonymous namespace)::indexing_backward_kernel...	0.00%	0.000us	0.00%	0.000us	0.000us	5.081s	2.43%	5.081s	714.669us	7110
void at::native::legacy::elementwise_kernel<128, 4, ...	0.00%	0.000us	0.00%	0.000us	0.000us	4.338s	2.07%	4.338s	220.297us	19690
aten::add_	0.12%	511.057ms	0.13%	565.776ms	11.515us	3.939s	1.88%	4.012s	81.656us	49132
aten::index	0.14%	630.846ms	8.33%	36.907s	1.647ms	3.616s	1.73%	6.082s	271.344us	22414
void at::native::index_elementwise_kernel<128, 4, vo...	0.00%	0.000us	0.00%	0.000us	0.000us	3.576s	1.71%	3.576s	204.712us	17470
void at::native::modern::elementwise_kernel<at::nati...	0.00%	0.000us	0.00%	0.000us	0.000us	3.006s	1.44%	3.006s	130.565us	23023
Cijk_Ailk_Bjlk_SB_MT128x64x16_SN_APM1_AF0EM1_AF1EM1_...	0.00%	0.000us	0.00%	0.000us	0.000us	2.842s	1.36%	2.842s	1.216ms	2338
MemcpyDeviceToDevice	0.00%	0.000us	0.00%	0.000us	0.000us	2.653s	1.27%	2.653s	84.007us	31583
aten::fill_	0.17%	774.034ms	0.18%	787.328ms	13.181us	2.646s	1.26%	2.690s	45.030us	59731
void at::native::modern::elementwise_kernel<at::nati...	0.00%	0.000us	0.00%	0.000us	0.000us	2.637s	1.26%	2.637s	51.976us	50736
Cijk_Alik_Bljk_SB_MT128x128x8_SN_APM1_AF0EM1_AF1EM1_...	0.00%	0.000us	0.00%	0.000us	0.000us	2.342s	1.12%	2.342s	988.252us	2370

RTX3090

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA %	CUDA total	CUDA time avg	# of Calls
aten::mm	1.42%	1.428s	1.58%	1.581s	32.469us	27.249s	40.50%	27.600s	566.772us	48696
aten::_index_put_impl_	-7.89%	-7910158.000us	42.18%	42.261s	1.972ms	11.035s	16.40%	13.147s	613.429us	21432
void (anonymous namespace)::indexing_backward_kernel...	0.00%	0.000us	0.00%	0.000us	0.000us	10.671s	15.86%	10.671s	675.386us	15800
void cutlass::Kernel<cutlass_80_tensorop_s1688gemm_1...	0.00%	0.000us	0.00%	0.000us	0.000us	10.047s	14.93%	10.047s	924.139us	10872
aten::mul	1.99%	1.997s	2.35%	2.350s	18.947us	8.920s	13.26%	9.544s	76.942us	124046
void cutlass::Kernel<cutlass_80_tensorop_s1688gemm_1...	0.00%	0.000us	0.00%	0.000us	0.000us	6.350s	9.44%	6.350s	803.806us	7900
void at::native::elementwise_kernel<128, 2, at::nati...	0.00%	0.000us	0.00%	0.000us	0.000us	6.015s	8.94%	6.015s	143.631us	41878
void cutlass::Kernel<cutlass_80_tensorop_s1688gemm_1...	0.00%	0.000us	0.00%	0.000us	0.000us	5.995s	8.91%	5.995s	843.171us	7110
cudaLaunchKernel	3.98%	3.991s	4.14%	4.149s	6.469us	5.449s	8.10%	5.542s	8.641us	641302
void at::native::vectorized_elementwise_kernel<4, at...	0.00%	0.000us	0.00%	0.000us	0.000us	4.644s	6.90%	4.644s	71.045us	65362
aten::copy_	0.93%	930.971ms	1.00%	1.093s	17.115us	3.885s	5.77%	4.031s	63.110us	63879
aten::add_	0.40%	486.193ms	0.62%	625.372ms	12.855us	3.220s	4.79%	3.389s	69.655us	48647
void at::native::vectorized_elementwise_kernel<4, at...	0.00%	0.000us	0.00%	0.000us	0.000us	2.529s	3.76%	2.529s	107.911us	23433
aten::fill_	0.56%	558.618ms	0.58%	585.615ms	9.804us	2.247s	3.34%	2.329s	38.990us	59733
void at::native::vectorized_elementwise_kernel<4, at...	0.00%	0.000us	0.00%	0.000us	0.000us	2.245s	3.34%	2.245s	44.264us	50708
aten::sum	0.43%	431.134ms	0.54%	541.348ms	24.049us	2.228s	3.31%	2.344s	104.133us	22510
void at::native::reduce_kernel<512, 1, at::native::R...	0.00%	0.000us	0.00%	0.000us	0.000us	2.204s	3.28%	2.204s	141.631us	15560
void at::native::elementwise_kernel<128, 2, at::nati...	0.00%	0.000us	0.00%	0.000us	0.000us	1.993s	2.96%	1.993s	89.349us	22305
Memcpy DtoD (Device -> Device)	0.00%	0.000us	0.00%	0.000us	0.000us	1.695s	2.52%	1.695s	97.585us	17367
aten::add	0.32%	318.336ms	0.38%	381.814ms	12.843us	1.358s	2.02%	1.475s	49.604us	29729
aten::index	-2.87%	-2874668.000us	13.62%	13.646s	608.837us	1.297s	1.93%	2.037s	90.884us	22414
void at::native::index_elementwise_kernel<128, 4, at...	0.00%	0.000us	0.00%	0.000us	0.000us	1.282s	1.91%	1.282s	73.371us	17470
fused_sigmoid_neg_add_mul_add_mul_mul	0.00%	0.000us	0.00%	0.000us	0.000us	1.169s	1.74%	1.169s	411.486us	2842
cudaPeekAtLastError	0.00%	1.232ms	0.00%	1.245ms	0.010us	1.130s	1.68%	1.130s	8.644us	130706
fused_mul_mul_mul_mu_24516804886151692 25	0.00%	0.000us	0.00%	0.000us	0.000us	1.102s	1.64%	1.102s	416.368us	2646
ampere_sgemm_64x64_nn	0.00%	0.000us	0.00%	0.000us	0.000us	1.062s	1.58%	1.062s	1.345ms	790
ampere_sgemm_32x128_tn	0.00%	0.000us	0.00%	0.000us	0.000us	728.994ms	1.08%	728.994ms	184.555us	3950
void cutlass::Kernel<cutlass_80_tensorop_s1688gemm_1...	0.00%	0.000us	0.00%	0.000us	0.000us	712.845ms	1.06%	712.845ms	300.778us	2370
cudaOccupancyMaxActiveBlocksPerMultiprocessorWithFla...	0.10%	101.388ms	0.10%	102.140ms	1.307us	670.390ms	1.00%	673.091ms	8.611us	78163

性能分析脚本

找到主函数，替换

显示底层计算库算子耗时

```
if __name__ == "__main__":
#     cProfile.run('main(running_as_script=True)')

    with torch.profiler.profile(
    activities=[
        torch.profiler.ProfilerActivity.CPU,
        torch.profiler.ProfilerActivity.CUDA,
    ]
) as p:
        main(running_as_script=True)
    print(p.key_averages().table(sort_by="self_cuda_time_total", row_limit=-1))
```

torchboard的分析

```
with torch.profiler.profile(
activities=[
torch.profiler.ProfilerActivity.CPU,
torch.profiler.ProfilerActivity.CUDA,
],
on_trace_ready=torch.profiler.tensorboard_trace_handler('./log/torchboard'),
record_shapes=True,
profile_memory=True,
with_stack=True) as p:
    main(running_as_script=True)
```

Magpy调研与应用，失败

后续方向

先重新保存模型 不用script model

尝试TVM编译model，Magpy